From scratch, single file implementation of a scaled-down version of MuZero for CartPole/Atari(WIP)

It's possible to implement key parts of famous algorithms at small scale from scratch in ~500 lines of code

AI Tinkerers Zurich, Jua.ai
April 9th, 2026

Adrian Scheerer
adrianscheerer@gmail.com

github.com/adrische/MuZero-MsPacman

linkedin.com/in/adrianscheerer/

Background & Motivation

Reinforcement learning:

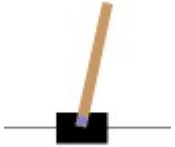

- State → Action → Reward → New state
- Maximize total expected discounted reward

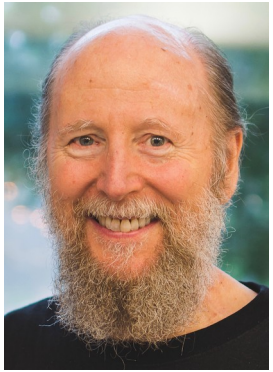

Bitter lesson for LLM pretraining?

incompleteideas.net/IncIdeas/BitterLesson.html

nature

Explore content | About the journal | Publish with us | Subscribe

nature > articles > article

Article | Published: 23 December 2020

Mastering Atari, Go, chess and shogi by planning with a learned model

Julian Schrittwieser, Ioannis Antonoglou, Thomas Hubert, Karen Simonyan, Laurent Sifre, Simon Schmitt, Arthur Guez, Edward Lockhart, Demis Hassabis, Thore Graepel, Timothy Lillicrap & David Silver

Nature **588**, 604–609 (2020) | Cite this article

66k Accesses | **1470** Citations | **1581** Altmetric | Metrics

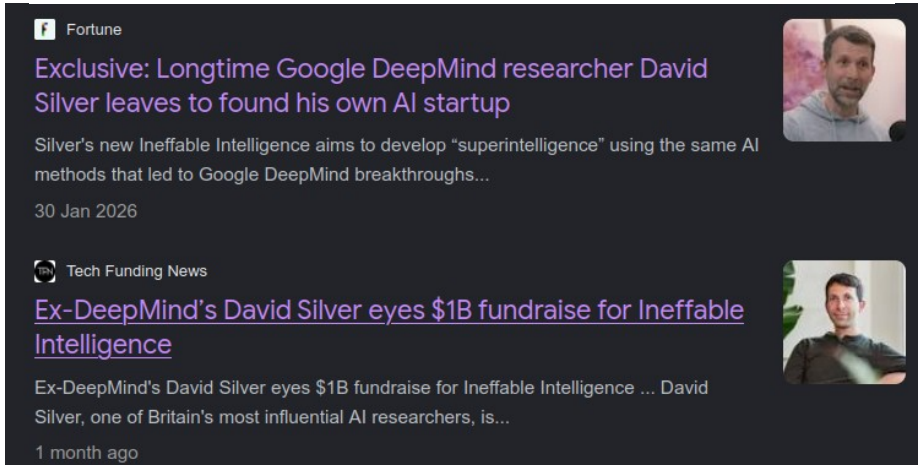

Fortune

Exclusive: Longtime Google DeepMind researcher David Silver leaves to found his own AI startup

Silver's new Ineffable Intelligence aims to develop "superintelligence" using the same AI methods that led to Google DeepMind breakthroughs...

30 Jan 2026

Tech Funding News

Ex-DeepMind's David Silver eyes $1B fundraise for Ineffable Intelligence

Ex-DeepMind's David Silver eyes $1B fundraise for Ineffable Intelligence ... David Silver, one of Britain's most influential AI researchers, is...

1 month ago